

Designing Intelligent Agents

EMATM0042 – Intelligent Information Systems

Monday 11 March – Part 2

kevin.mcareahey@bristol.ac.uk

Previous Lecture...

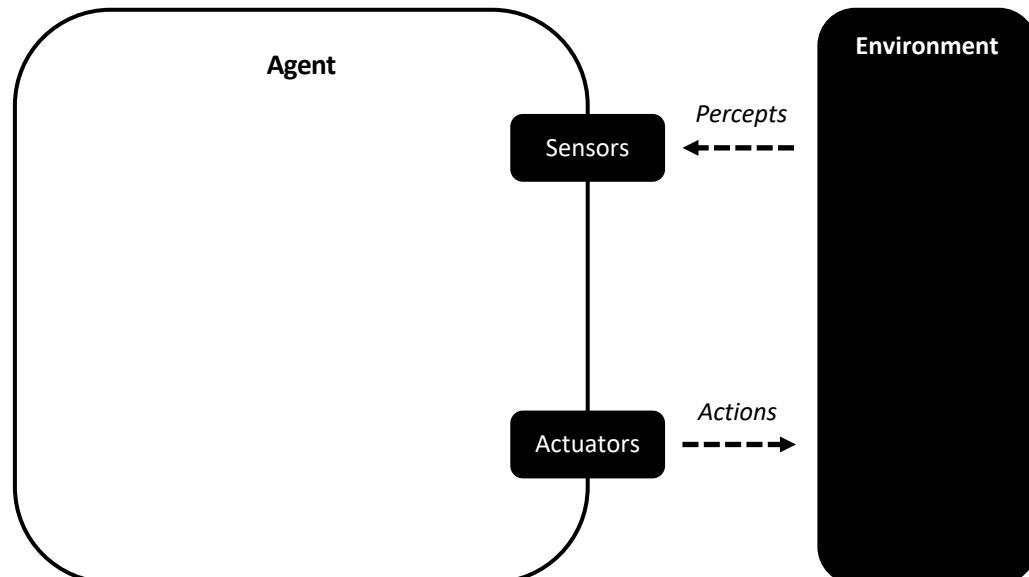
1. Background
 - Trends in computing
2. Informal definition
 - Agent
 - MAS
3. Understanding the problem
 - Micro vs. macro
 - Engineering vs. theory
4. Applications

This Lecture...

1. Abstract agent representations
 - Agent function
 - Agent program
2. Task environments
 - Problem Analysis
 - Properties
3. Types of agents
 - Reactive agents
 - Proactive agents
 - Cognitive agents

Agents

Basic Definition

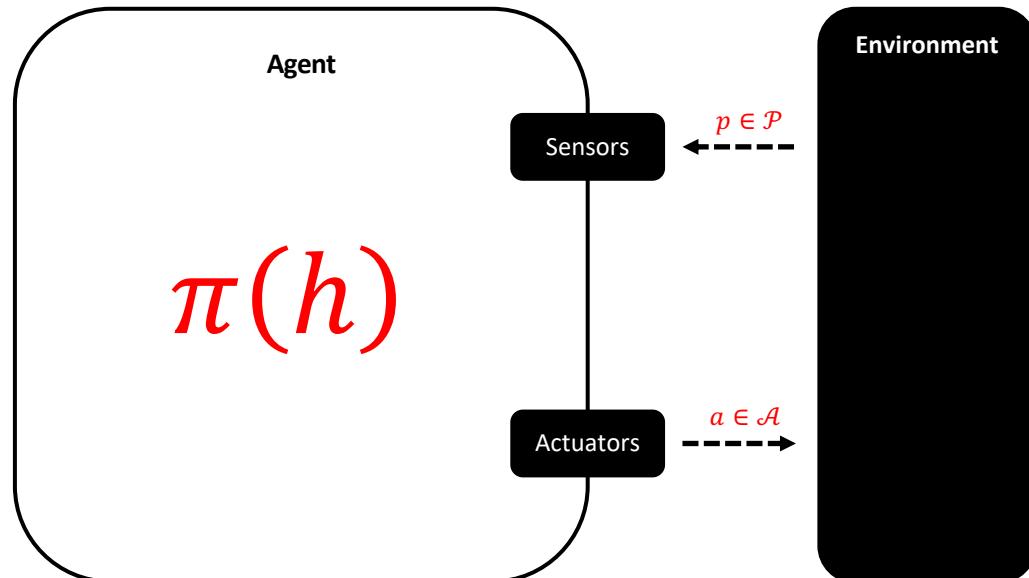


- An **agent** is an entity that **acts**
 - Situated in an **environment**
 - Acts upon its environment through **actuators**
 - Perceives its environment though **sensors**

Agent	Environment	Actuators	Sensors
Human?	Real-world	Hands, vocals, etc.	Eyes, ears, etc.
Robot?	Real-world	Motors, etc.	Cameras, etc.
Internet bot?	Internet	HTTP POST, etc.	Text search, etc.
Light switch?	Circuit	Current	Switch

Agents

Basic Definition: Formal

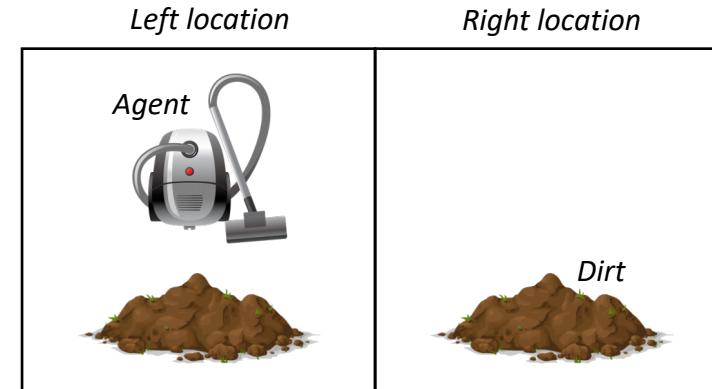


- Set of **actions** $\mathcal{A} = \{a, a', \dots\}$
 - Set of **percepts** $\mathcal{P} = \{p, p', \dots\}$
-] Atomic symbols
-
- A **history** is a finite sequence $h = (a, p, \dots, a', p')$
 - Action-percept sequence starting with an action and ending with a percept
 - Complete record of what has happened in the agent's life
 - Set of **histories** $H = \{h, h', \dots\}$
 - An **agent function** is a function $\pi : H \rightarrow \mathcal{A}$
 - A mapping from histories to actions

Agents

Basic Definition: Formal – Example

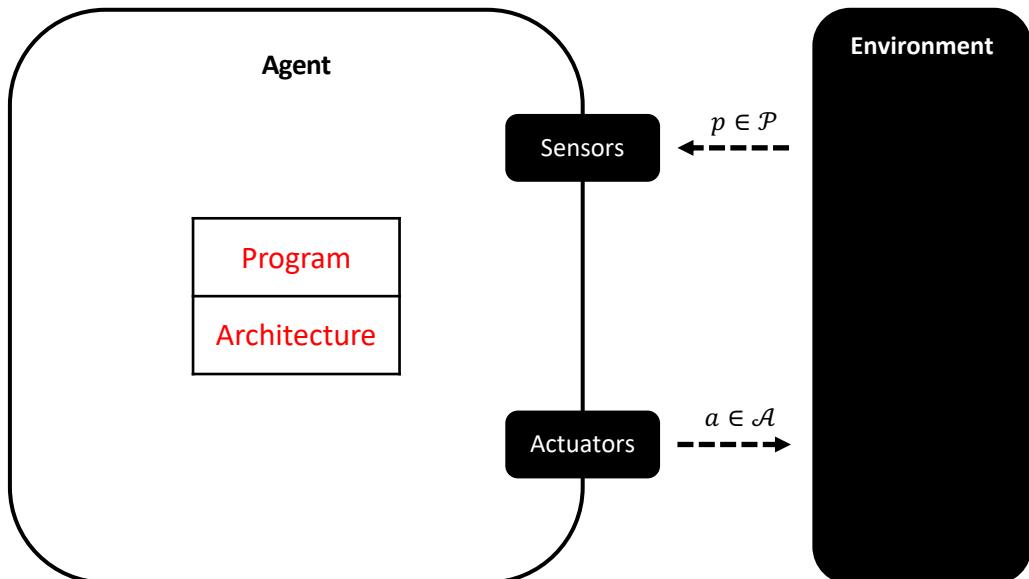
- Suppose we have a vacuum-cleaner agent...
 - Assume that it **persists** forever
 - Assume that its sensing is restricted to its **current location**
- Actions $a \in \mathcal{A}$
 - **Move(Right)**: move to right location
 - **Move(Left)**: move to left location
 - **Clean**: remove dirt from current location
- Percepts $p \in \mathcal{P}$
 - **Sensed(Left, Dirty)**: dirt in left location
 - **Sensed(Left, Clean)**: no dirt in left location
 - **Sensed(Right, Dirty)**: dirt in right location
 - **Sensed(Right, Clean)**: no dirt in right location



History $h \in H$	Agent function $\pi(h)$
()	Clean
(Clean, Sensed(Left, Clean))	Move(Right)
(Clean, Sensed(Right, Clean))	Move(Left)
...	...
(Clean, Sensed(Left, Clean), Move(Right), Sensed(Right, Dirty))	Clean
(Clean, Sensed(Left, Clean), Move(Right), Sensed(Right, Clean))	Move(Left)
...	...

Agents

Implementation



- Agent function is just an **abstract** mathematical description of an agent
 - Set of histories is very **large or infinite**
 - Says nothing about how to **choose** action $\pi(h)$
- Need a way to **implicitly** define an agent function
 - **Concise** representation
 - **Decide** what action to perform
- An **agent program** implements an agent function
 - This is our main focus
- An **agent program** runs on an **agent architecture**
 - Physical robot, virtual machine, agent interpreter, etc.

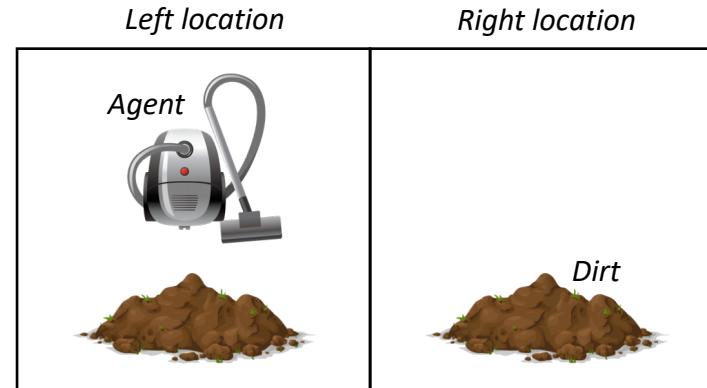
Agents

Implementation – Example

Algorithm: Vacuum-cleaner agent program

```
1 EXECUTE(Clean)
2 while true do
3   Sensed( $x, y$ )  $\leftarrow$  SENSE()
4   if  $y = \text{Dirty}$  then
5     EXECUTE(Clean)
6   else if  $x = \text{Left}$  then
7     EXECUTE(Move(Right))
8   else
9     EXECUTE(Move(Left))
```

Implicitly defines...



History $h \in H$	Agent function $\pi(h)$
()	Clean
(Clean, Sensed(Left, Clean))	Move(Right)
(Clean, Sensed(Right, Clean))	Move(Left)
...	...
(Clean, Sensed(Left, Clean), Move(Right), Sensed(Right, Dirty))	Clean
(Clean, Sensed(Left, Clean), Move(Right), Sensed(Right, Clean))	Move(Left)
...	...

Agents

Intelligent Agents

- An agent is **autonomous** if it does not rely exclusively on prior knowledge
 - The agent should at least account for its own **experience**
- An agent is **rational** if it acts towards achieving the “best” outcome
 - Behaviour should fit some criterion for **success**
- We are typically interested in agents that are **autonomous and rational**
 - We might call these **intelligent agents**, but the name is not important...



Agents

Understanding the Problem

- An agent is a **complete** system that solves a **specific problem** by integrating a range of (often shallow) competences
- Analysis of the **task environment** helps us understand what competences are important and what are not

Task Environment = (Environment, Actuators, Sensors, Performance Measure)

What is **relevant** to describing the environment that the agent will **inhabit**?

What actuators are **available** to the agent and how do they **affect** the environment?

What sensors are **available** to the agent and what **feedback** do they provide about the environment?

What is our measure of **success** for the agent?

Agents

Understanding the Problem – Example

Agent	Environment	Actuators	Sensors	Performance Measure
Driverless car	Roads, other vehicles, pedestrians, weather	Steering, accelerator, brake, signal, horn, display	Cameras, sonar, speedometer, GPS, odometer, accelerometer, engine sensors, keyboard	Safety, speed, legality, comfort
Medical diagnosis system	Patient, hospital, staff	Display of questions, tests, diagnoses, treatments, referrals	Keyboard entry of symptoms, findings, patient's answers	Healthy patient, reduced costs
Satellite image analysis system	Downlink from orbiting satellite	Display of scene categorization	Color pixel arrays	Correct image categorization
Part-picking robot	Conveyor belt with parts; bins	Jointed arm and hand	Camera, joint angle sensors	Percentage of parts in correct bins
Refinery controller	Refinery, operators	Valves, pumps, heaters, displays	Temperature, pressure, chemical sensors	Purity, yield, safety
Interactive English tutor	Set of students, testing agency	Display of exercises, suggestions, corrections	Keyboard entry	Student's score on test

Task Environment

Properties

- Deciding how to **model** a task environment is crucial to choosing an **appropriate** agent design
 - What **properties** does it exhibit?
 - What design **trade-offs** are appropriate?
- In practice, we need to impose simplifying **assumptions** or **restrictions** in our model
 - Only possible to **approximate** many real task environments (e.g. robots in the real-world)
 - **Computational** considerations
- Allow more complexity for properties that have greater impact on the **performance measure**, and less complexity for others



#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Unknown

Task Environment

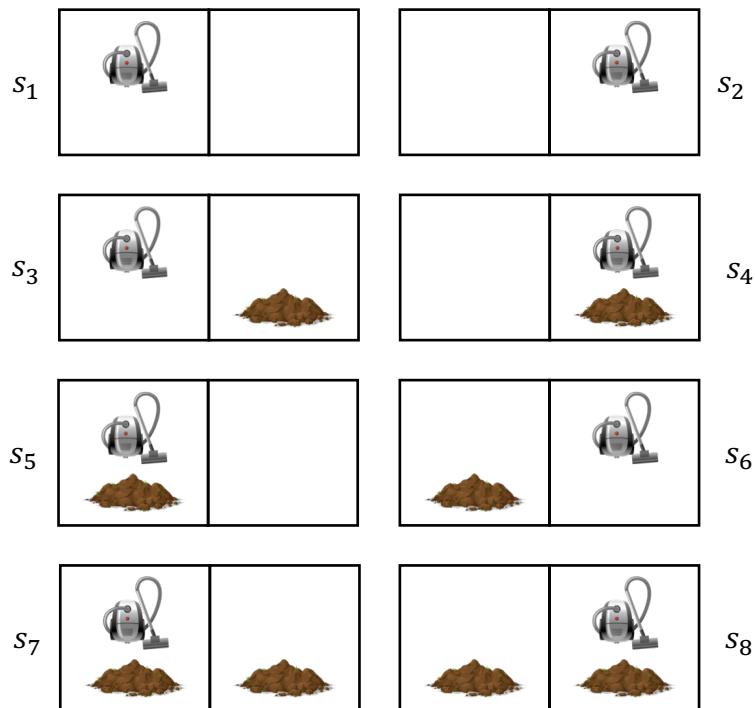
Properties: (1) Discrete vs. Continuous

- ...**discrete** if there is a **finite** number of **environment states, actions, and percepts**
- ...**continuous** if there is **not a finite** number of environment states, actions, or percepts
 - Subsumes discrete task environments (e.g. continuous environment states with discrete actions and percepts)

#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Unknown

Task Environment

Properties: (1) Discrete vs. Continuous – Example



- Discrete environment states $\mathcal{S} = \{s_1, s_2, \dots, s_8\}$
 - Suppose we include (real) time in the environment states?
- Discrete actions $\mathcal{A} = \left\{ \begin{array}{l} \text{Move(Left),} \\ \text{Move(Right),} \\ \text{Clean} \end{array} \right\}$
- Discrete percepts $\mathcal{P} = \left\{ \begin{array}{l} \text{Sensed(Left, Clean),} \\ \text{Sensed(Left, Dirty),} \\ \text{Sensed(Right, Clean),} \\ \text{Sensed(Right, Dirty)} \end{array} \right\}$

...this environment is **discrete**

Task Environment

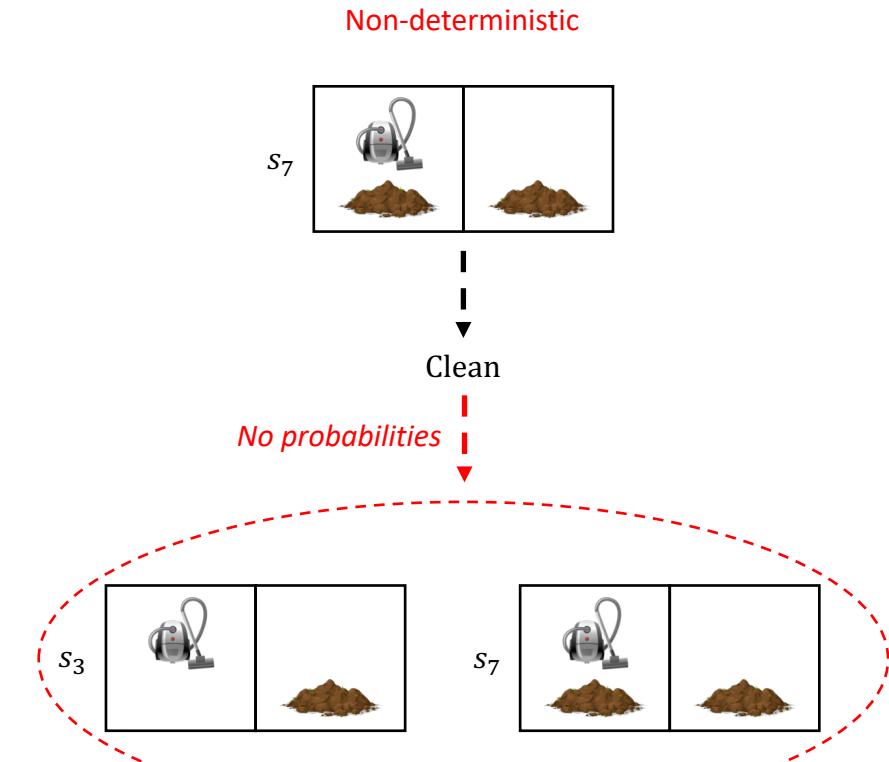
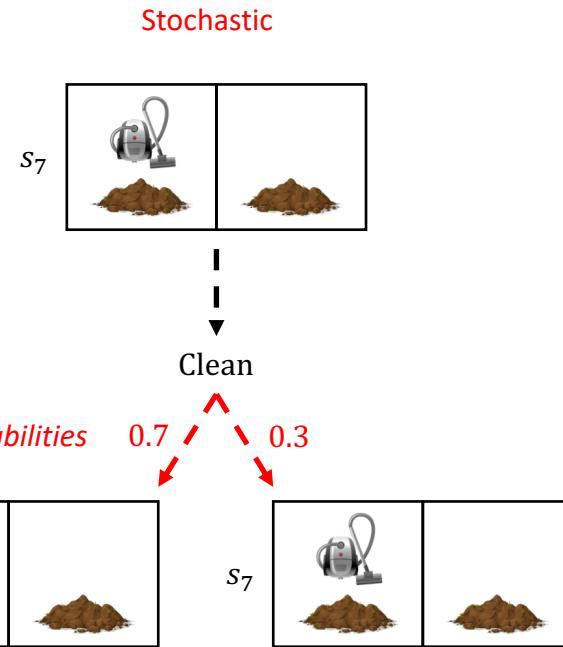
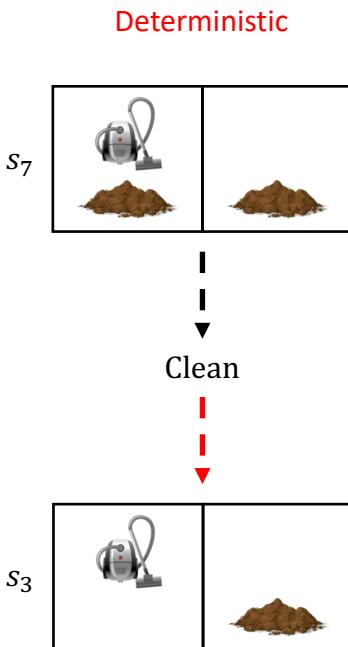
Properties: (2) Deterministic vs. Stochastic

- ...**deterministic** if each action (or event) leads from one environment state to a **unique** successor state
- ...**non-deterministic** if each action (or event) leads to **one-of-many** possible successor states
 - Subsumes deterministic task environments
- ...**stochastic** if it is non-deterministic and there are **probabilities** associated with successor states
 - Subsumes non-deterministic task environments, assuming that possible successors states have equal (non-zero) probability of occurring

#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Unknown

Task Environment

Properties: (2) Deterministic vs. Stochastic – Example



Task Environment

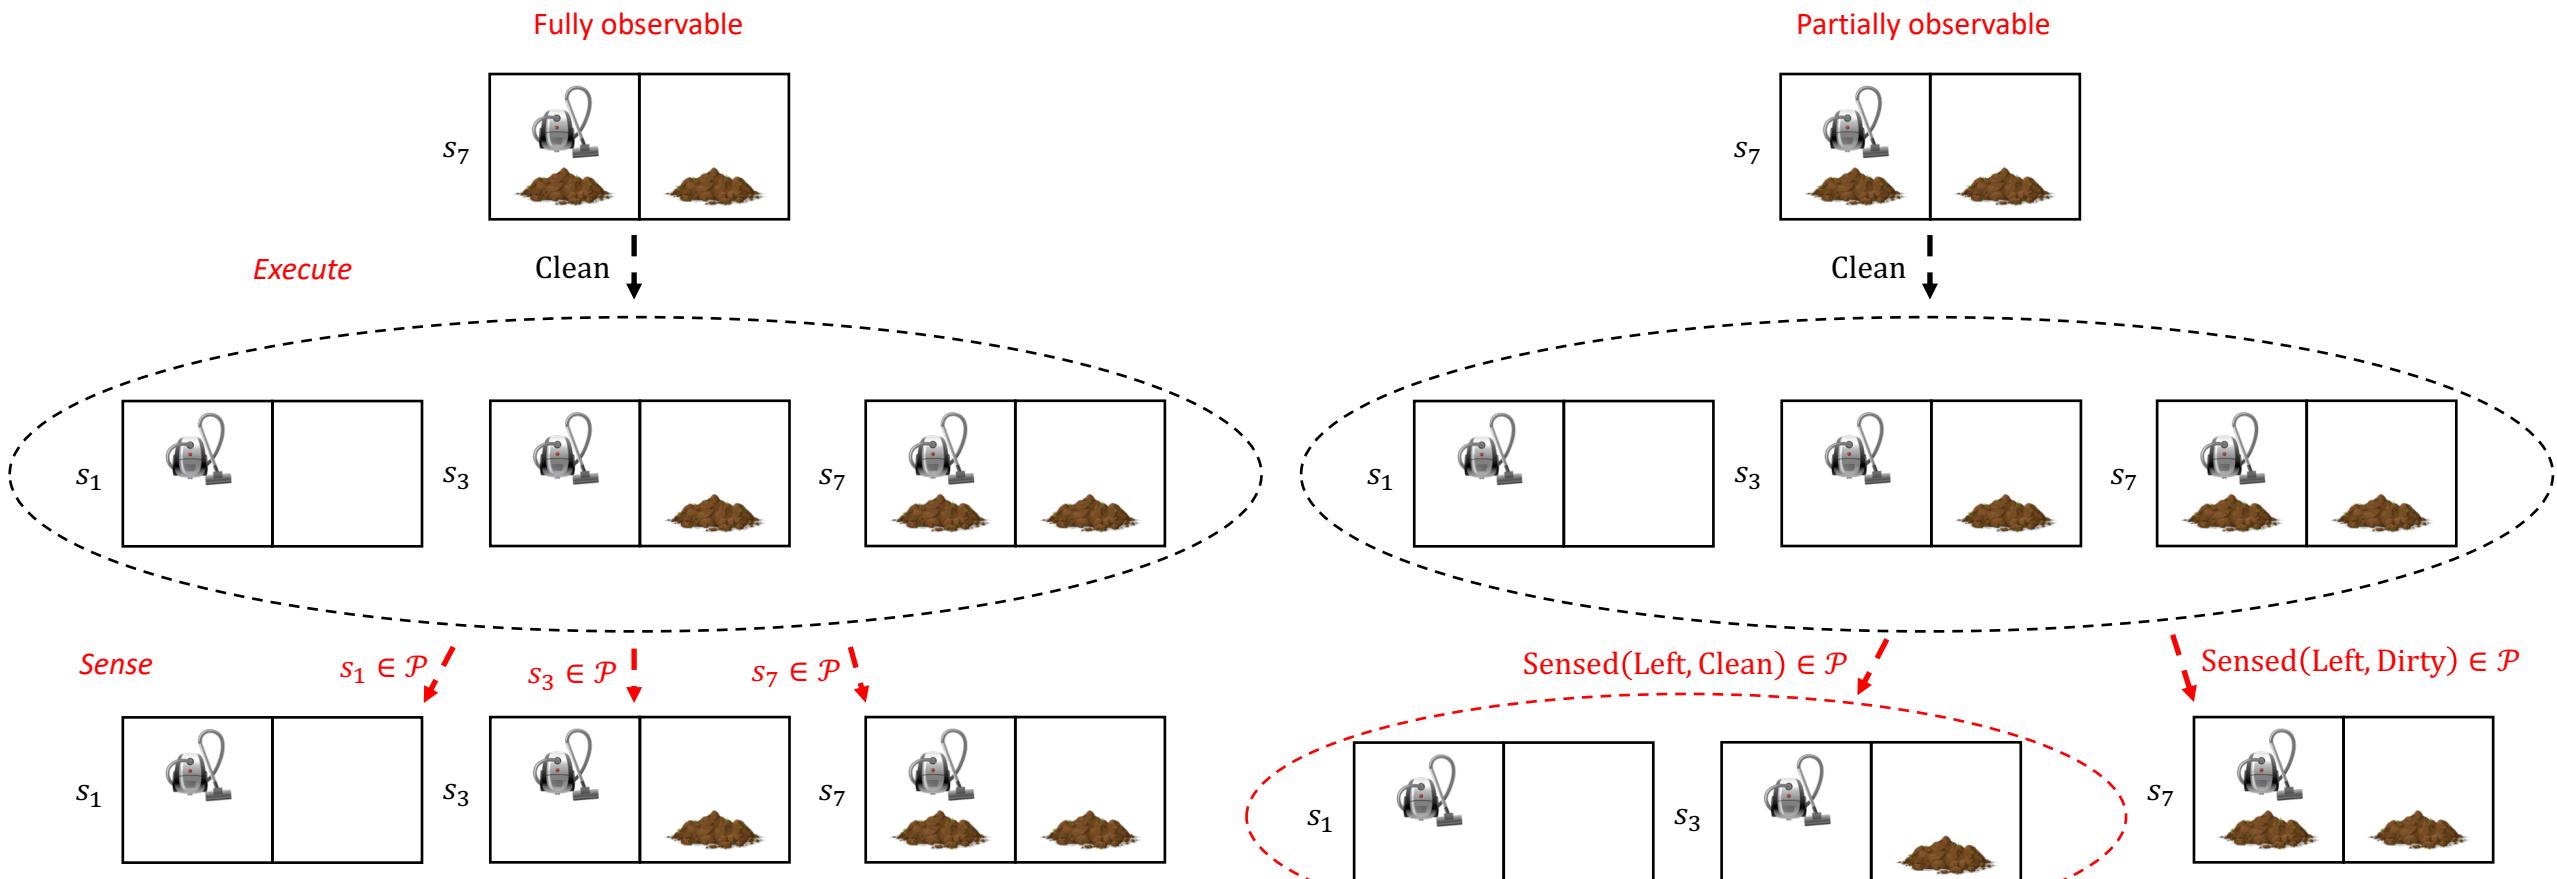
Properties: (3) Fully Observable vs. Partially Observable

- ...**fully observable** if percepts provide **complete information** about current environment state
 - Agent **receives feedback** and is always **certain** about its current environment state when executing actions
- ...**partially observable** if percepts provide **incomplete information**
 - Agent **receives feedback** but may be **uncertain** about its current environment state when executing actions
 - Subsumes fully observable environments
- ...**non-observable** if percepts provide **no information**
 - Agent **receives no feedback** and may be **uncertain** about its current environment state when executing actions
 - Subsumed by partially observable environments

#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Unknown

Task Environment

Properties: (3) Fully Observable vs. Partially Observable – Example



Task Environment

Properties: (4) Episodic vs. Sequential

- ...**episodic** if the agent's decisions occur in discrete episodes with **no dependencies** on the past or future
 - Agent chooses actions based on its **latest percept** only, and **does not consider** long-term consequences
- ...**sequential** if the agent's decisions **depend** on its past and/or its expectations about the future
 - Agent chooses actions based on its **history**, and considers the long-term **consequences** of its actions
 - Subsumes episodic task environments

#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Unknown

Task Environment

Properties: (4) Episodic vs. Sequential – Example



Episodic: part-picking robot



Sequential: chess

Task Environment

Properties: (5) Static vs. Dynamic

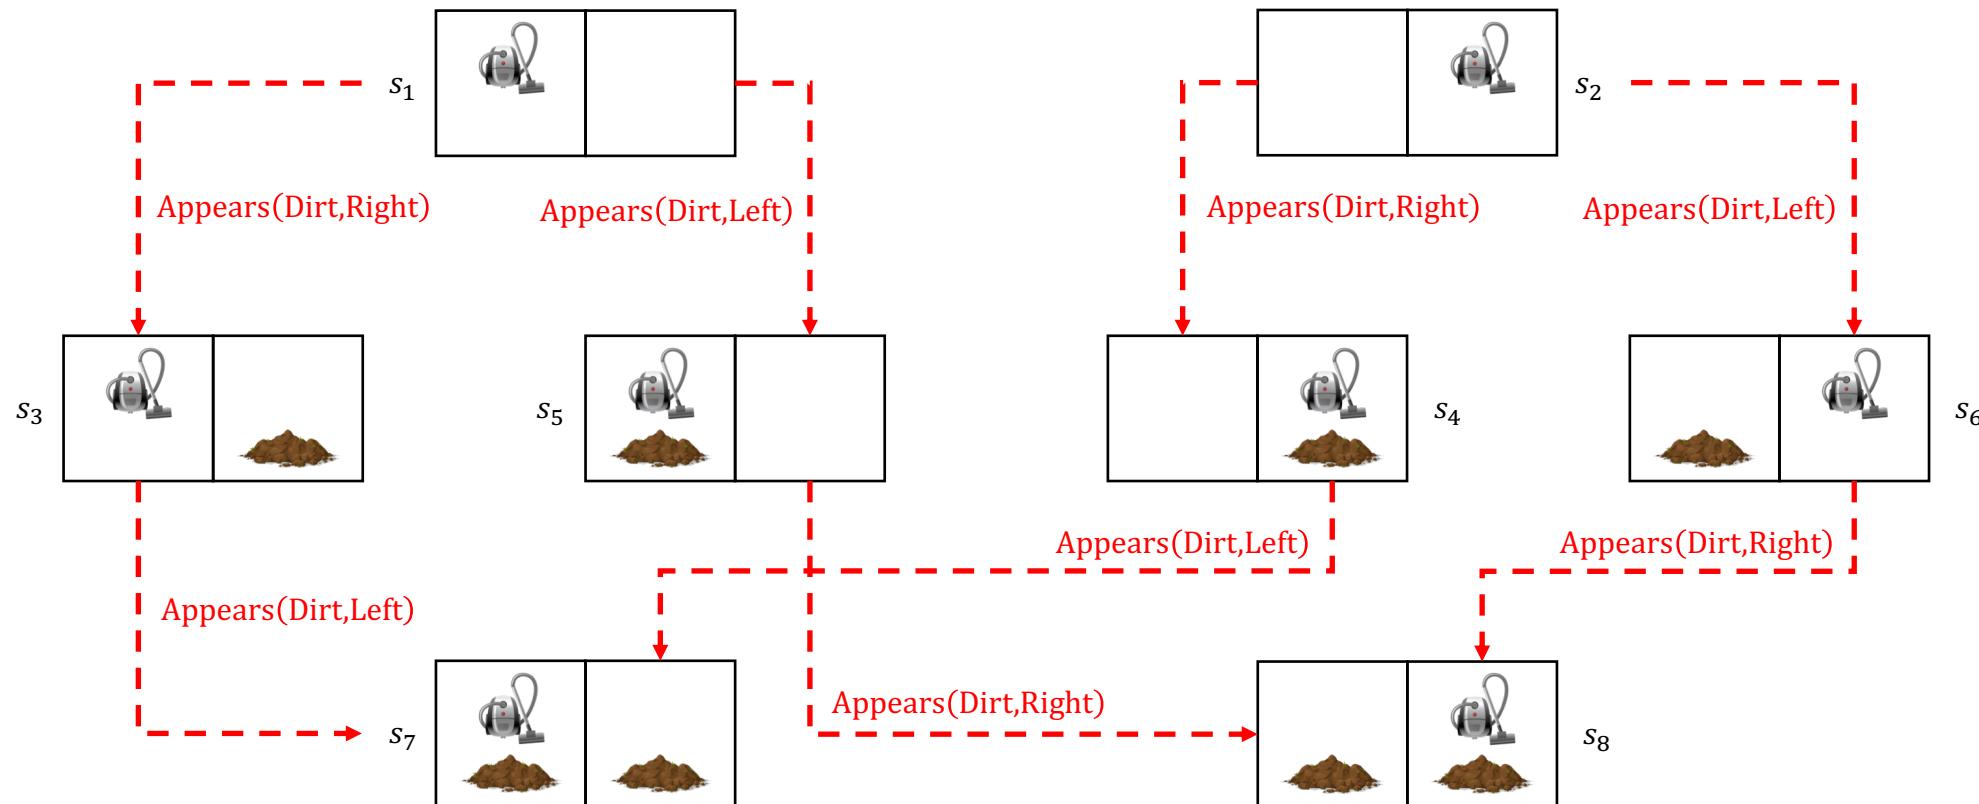
- ...**static** if there are **no exogenous events**
 - Agent has complete control over the environment
- ...**dynamic** if there may be **exogenous events**
 - Environment state may change due to external events

#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Unknown

Task Environment

Properties: (5) Static vs. Dynamic – Example

Exogenous events $\mathcal{E} = \{\text{Appears(Dirt,Left)}, \text{Appears(Dirt,Right)}\}$



Task Environment

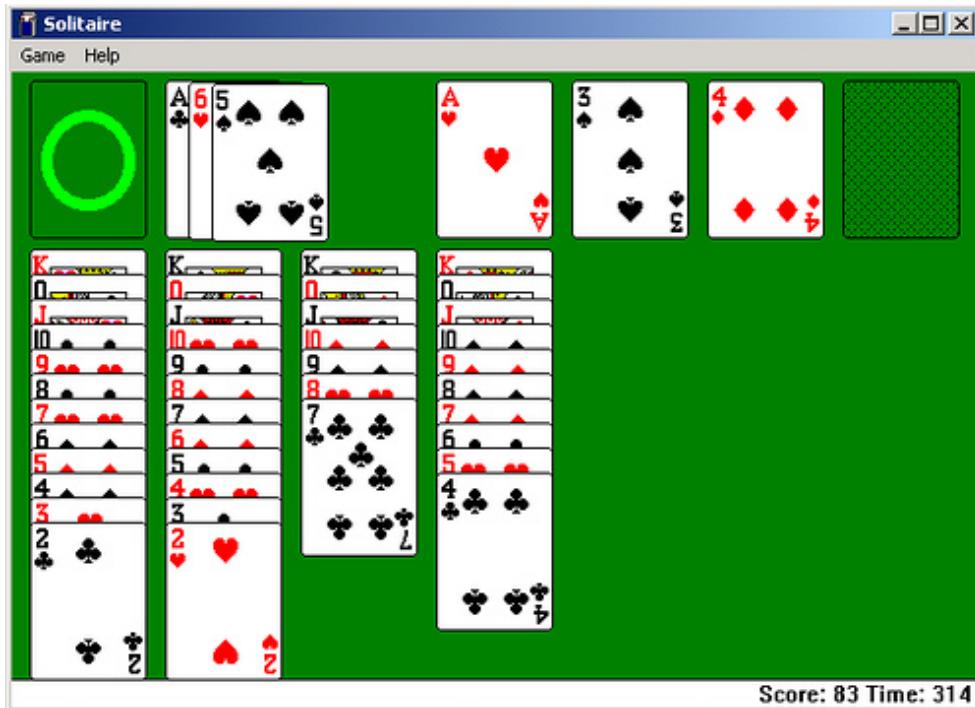
Properties: (6) Single Agent vs. Multi-Agent

- ...**single agent** if the agent does not account for the behaviour other agents
 - Might be the only agent that inhabits the environment
 - Might treat the actions of other agents as exogenous events (i.e. single agent but dynamic)
- ...**multi-agent** if the agent accounts for the behavior of other agents when choosing its own actions
 - Cooperation, competition, etc.
 - Subsumes single agent task environments

#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Unknown

Task Environment

Properties: (6) Single Agent vs. Multi-Agent – Example



Single agent: solitaire



Multi-agent: poker

Task Environment

Properties: (7) Known vs. Unknown

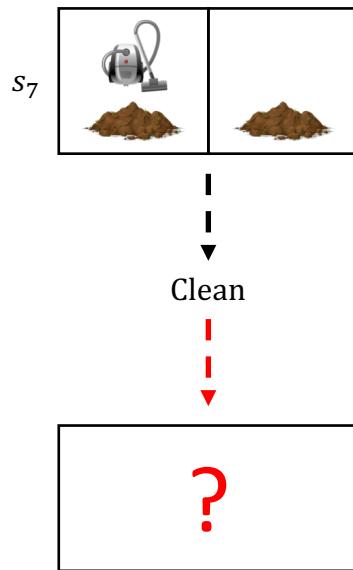
- ...**known** if the agent has a **complete model** of the task environment
 - Model may be imperfect, but is typically fixed
- ...**unknown** if the agent has an **incomplete model** (or no model) of the task environment
 - Agent may be required to learn (or improve) the model

#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Unknown

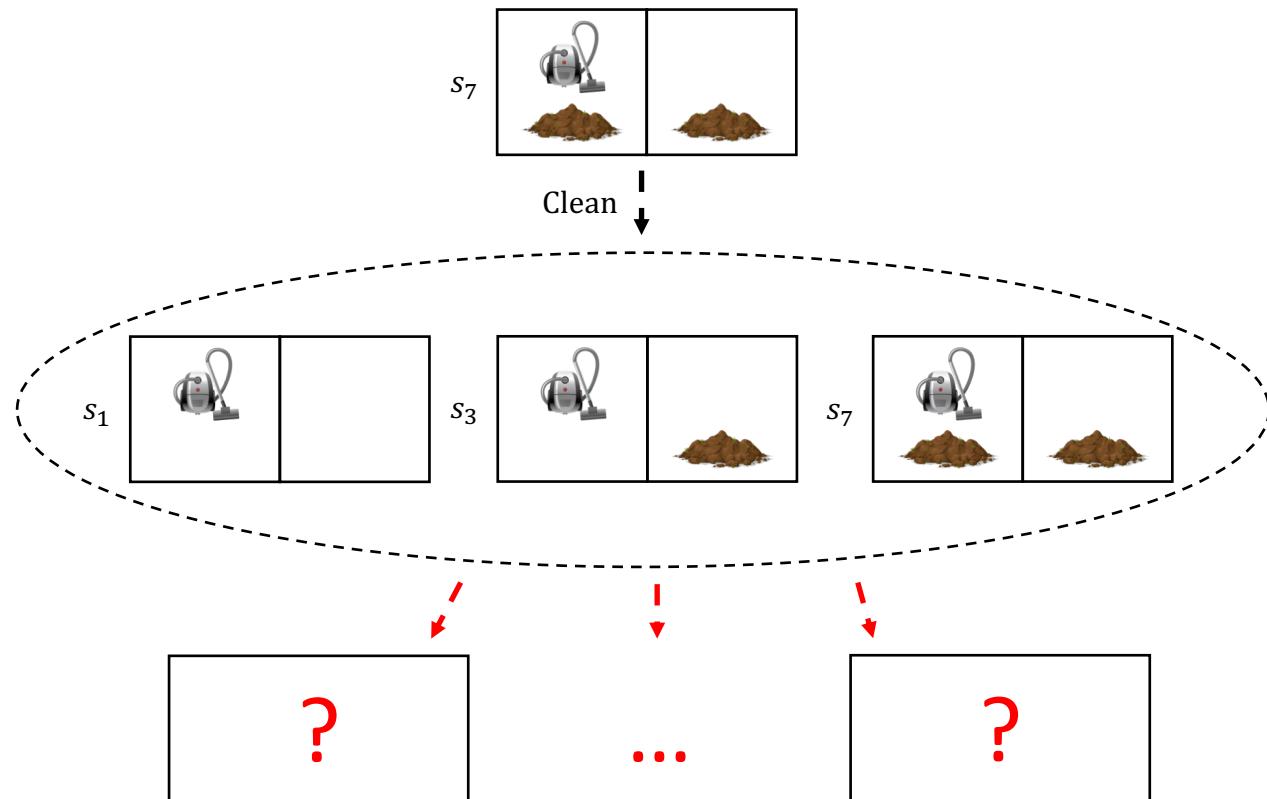
Task Environment

Properties: (7) Known vs. Unknown – Example

Unknown actuator model



Unknown sensor model



Types of Agents

Overview

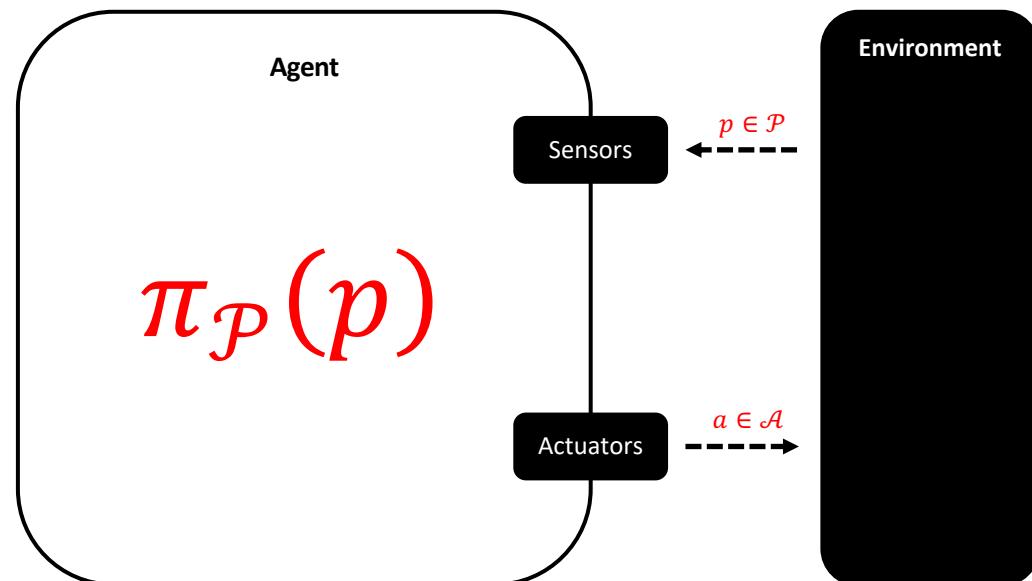
- Reactive agents
 - Memoryless
 - With belief state
- Proactive agents
 - Goal-based agents
 - Utility-based agents
- Cognitive agents
 - Belief-desire-intention agents

We will touch on each of these properties...

#	Less Challenging	More Challenging
1	Discrete	Continuous
2	Deterministic	Stochastic (or non-deterministic)
3	Fully observable	Partially observable (or non-observable)
4	Episodic	Sequential
5	Static	Dynamic
6	Single agent	Multi-agent
7	Known	Partially known (or unknown)

Reactive Agents

Memoryless



- A **memoryless** agent function is $\pi_{\mathcal{P}} : \mathcal{P} \rightarrow \mathcal{A}$
 - A mapping from percepts to actions
- Decides what to do based only on **latest percept**
 - Ignores rest of history
- Induces an agent function π
 - $\pi(h) = \pi_{\mathcal{P}}(p)$ for any $h \in H$ such that h ends with p
- Can work well when the agent has **full observability**
 - **Problematic** when the agent only has partial observability...

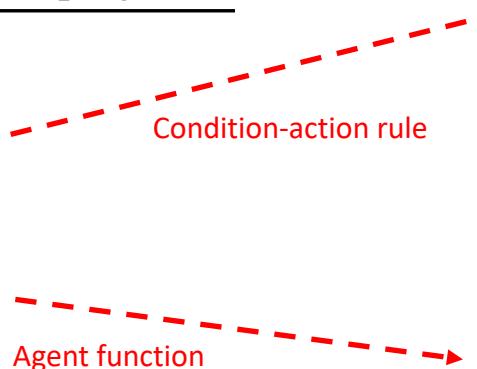
Reactive Agents

Memoryless – Example

Algorithm: Memoryless rule-based agent program

Input: Condition-action rule base K

```
1 while true do
2   p ← SENSE()
3   for each  $(\phi, a) \in K$  do
4     if  $p \models \phi$  then
5       EXECUTE( $a$ )
6       break
```

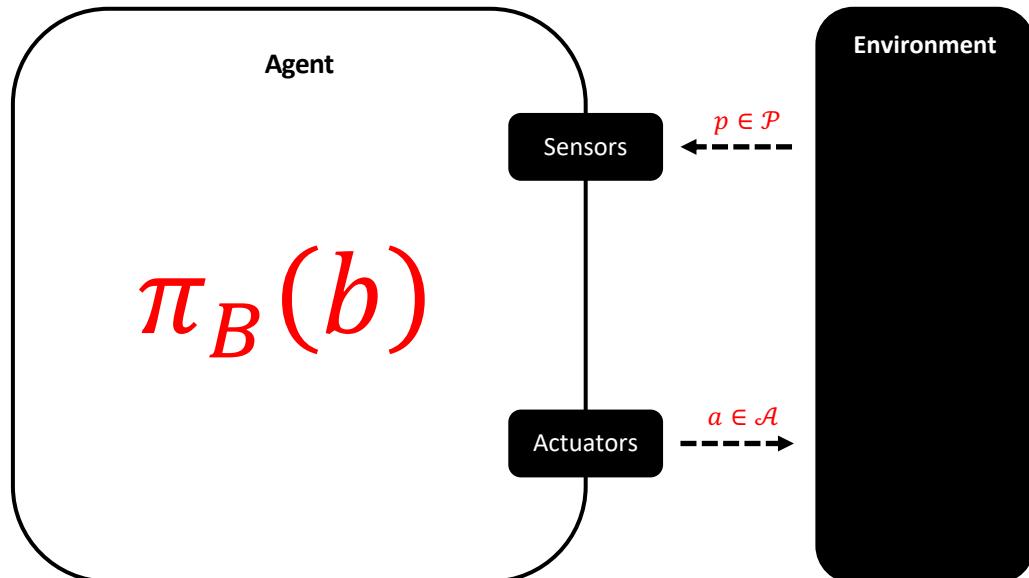


#	Condition ϕ	Action $a \in \mathcal{A}$
1	Sensed(x , Dirty)	Clean
2	Sensed(Left, Clean)	Move(Right)
3	Sensed(Right, Clean)	Move(Left)

Percept $p \in \mathcal{P}$	Memoryless agent function $\pi_{\mathcal{P}}(p)$
Sensed(Left, Dirty)	Clean
Sensed(Right, Dirty)	Clean
Sensed(Left, Clean)	Move(Right)
Sensed(Right, Clean)	Move(Left)

Reactive Agents

With Belief State



- Set of **belief states** B
 - A belief state $b \in B$ represents the agent's (potentially uncertain) view of the current environment state
- A **belief-based** agent function $\pi_B : B \rightarrow \mathcal{A}$
 - A mapping from belief states to actions
- Induces an agent function π
 - $\pi(h) = \pi_B(b)$ for any $h \in H$ such that b is the belief state associated with h
- How to **model** an agent's belief state?

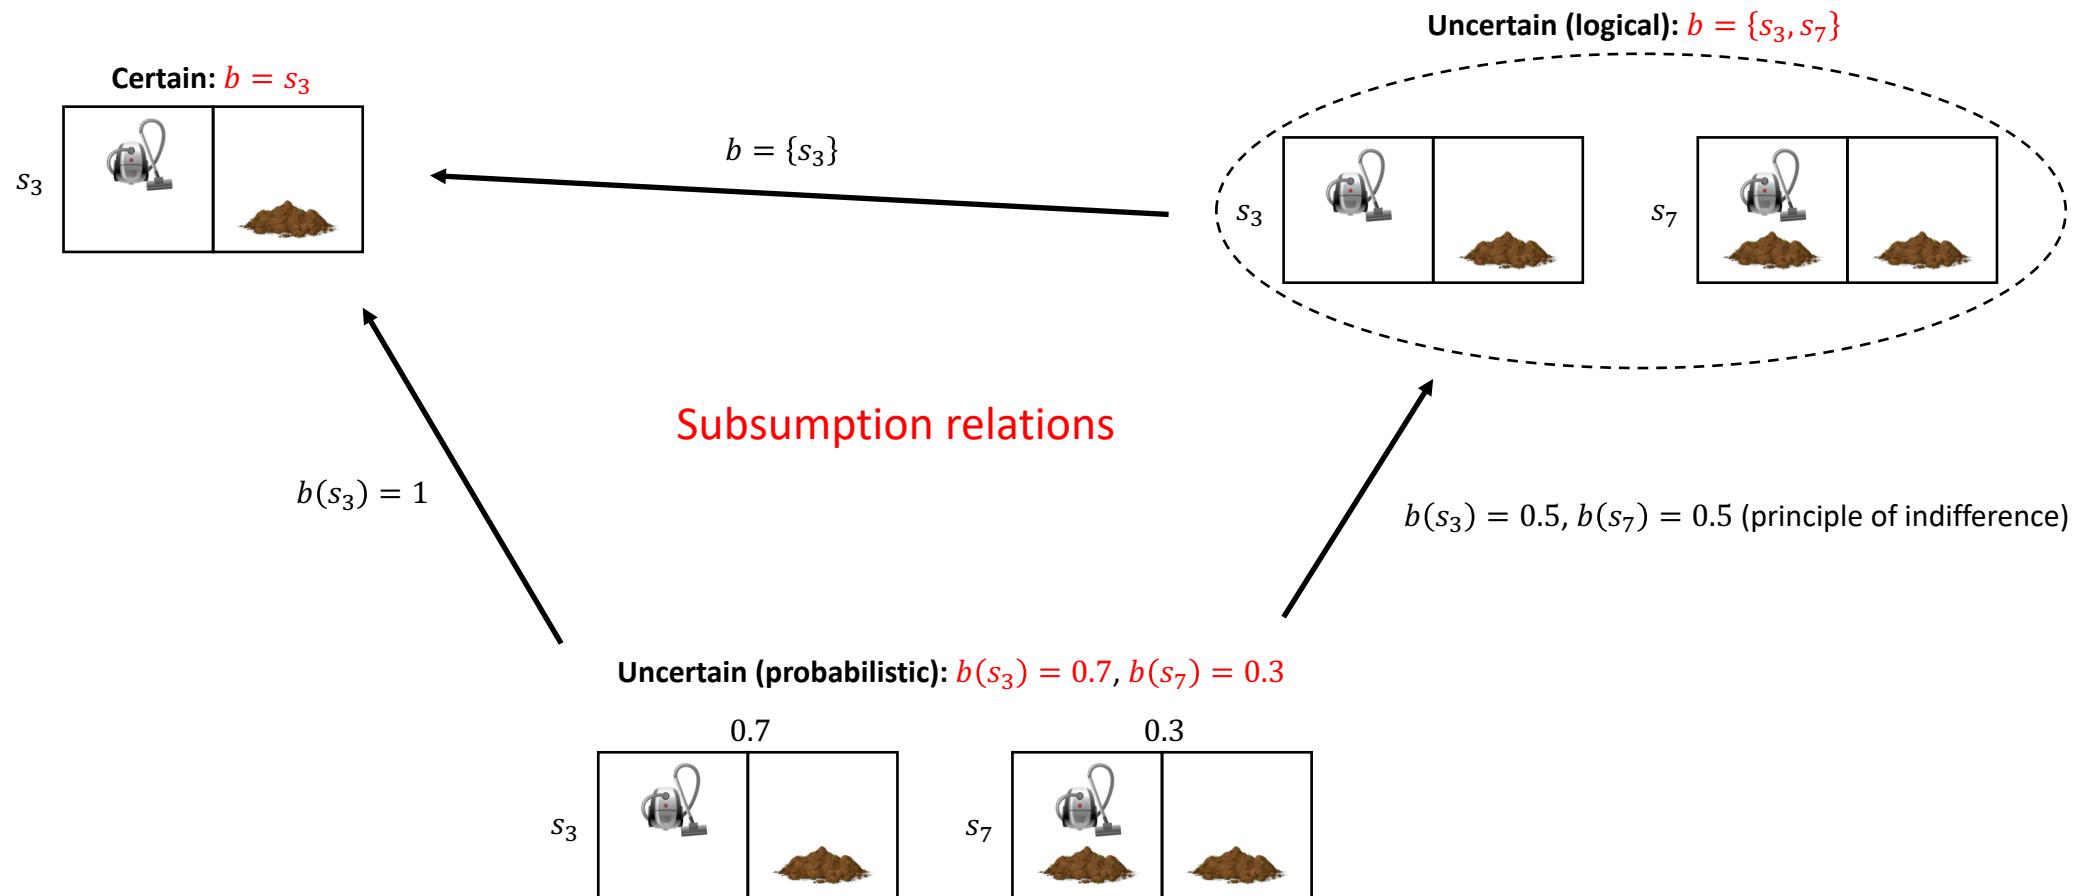
Reactive Agents

With Belief State: Belief State Model

- If the agent has **full observability**, then $B = S$
 - No uncertainty over the current environment state, since each percept uniquely identifies an environment state
- If the agent has only **partial observability**, then $B \neq S$, e.g.
 - $B = 2^S$ if we use a **logical** model of uncertainty
 - $B = \Pi(S)$ if we use a **probabilistic** model of uncertainty
- With partial observability, it is necessary to **track** the agent's current belief state
 - With full observability, a **memoryless agent** is sufficient
- How to **update** the belief state after each action-percept cycle?
 - Need an **actuator** model and a **sensor** model

Reactive Agents

With Belief State: Belief State Model – Example



Reactive Agents

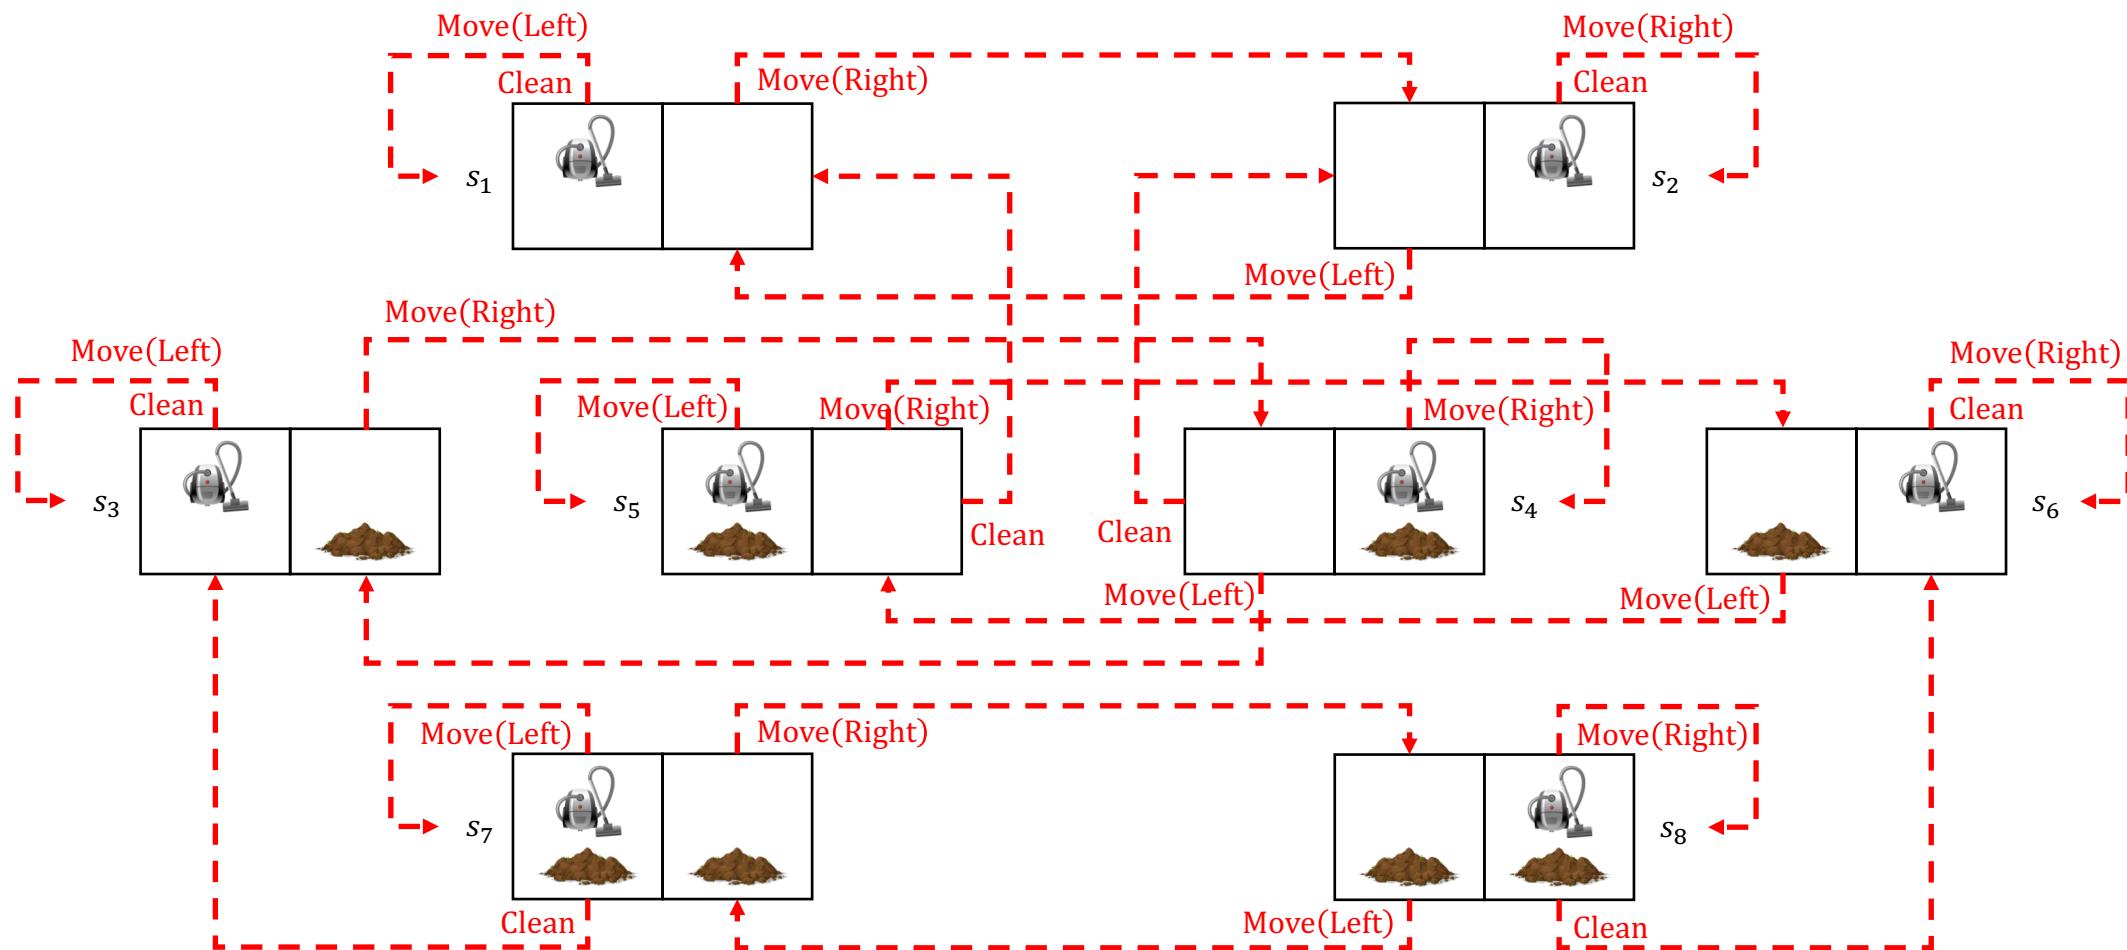
With Belief State: Actuator Model

- Actuator model is typically formalized as a **transition function** $T : \mathcal{S} \times \mathcal{A} \rightarrow \dots$
 - Effect of action depends only on **current environment state** (see e.g. Markov assumption)
- **Deterministic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$
 - Executing action $a \in \mathcal{A}$ in environment state $s \in \mathcal{S}$ results in a **single** successor state $T(s, a) \in \mathcal{S}$
- **Non-deterministic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow 2^{\mathcal{S}}$
 - Executing action $a \in \mathcal{A}$ in environment state $s \in \mathcal{S}$ results in a **set** of possible successor states $T(s, a) \subseteq \mathcal{S}$
- **Stochastic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \Pi(\mathcal{S})$
 - Executing action $a \in \mathcal{A}$ in environment state $s \in \mathcal{S}$ results in a **probability distribution** over successor states $T(s, a) \in \Pi(\mathcal{S})$

Reactive Agents

With Belief State: Actuator Model – Example

...this is a **deterministic** transition function



Reactive Agents

With Belief State: Sensor Model

- Sensor model is typically formalized as an **observation function** $O : \mathcal{S} \rightarrow \mathcal{P}$
 - If the agent arrives in environment state $s \in \mathcal{S}$, they will receive a **single** percept $O(s) \in \mathcal{P}$
 - May also **depend on action** as $O : \mathcal{A} \times \mathcal{S} \rightarrow \mathcal{P}$ (i.e. arriving in state $s \in \mathcal{S}$ after executing action $a \in \mathcal{A}$...)
- **Observability** depends on how O is defined
 - **Always** (weakly) partially observable
 - Fully observable if O is a **bijection** (i.e. a one-to-one mapping) from \mathcal{S} to \mathcal{P}
 - Non-observable if $O(s_1) = \dots = O(s_n) = p$ (i.e. p is meaningless)
 - **Strictly** partially observable if **neither** fully observable or non-observable
- The above is a **deterministic** observation function, but it is also possible to define a **noisy** sensing model
 - **Non-deterministic** observation function $O : \mathcal{S} \rightarrow 2^{\mathcal{P}}$
 - **Stochastic** observation function $O : \mathcal{S} \rightarrow \Pi(\mathcal{P})$

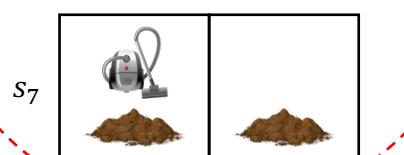
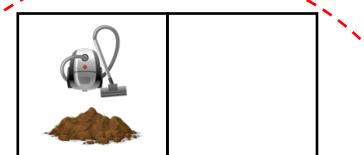
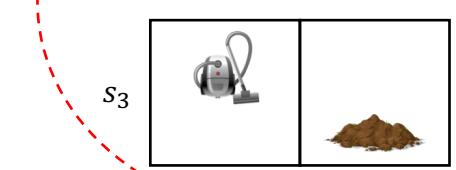
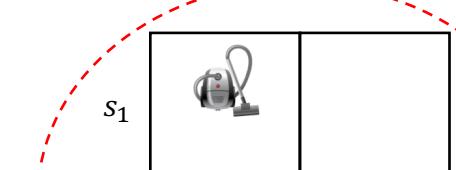
Reactive Agents

With Belief State: Sensor Model – Example

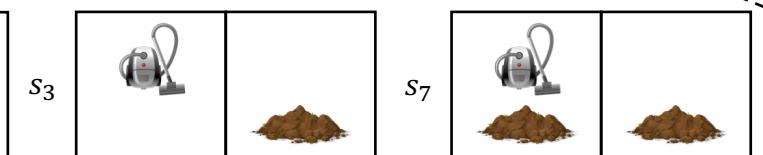
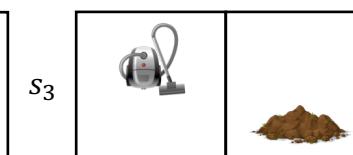
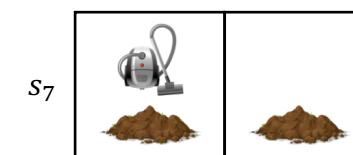
...this is a **deterministic** but **partially observable** observation function

$$O(s_1) = O(s_3) = \text{Sensed(Left, Clean)}$$

$$O(s_2) = O(s_6) = \text{Sensed(Right, Clean)}$$



Observation function



$$p \in \{O(s_1), O(s_3), O(s_7)\} \Leftrightarrow p \in \{\text{Sensed(Left, Clean)}, \text{Sensed(Left, Dirty)}\}$$

$$O(s_5) = O(s_7) = \text{Sensed(Left, Dirty)}$$

$$O(s_4) = O(s_8) = \text{Sensed(Right, Dirty)}$$

Reactive Agents

With Belief State: Update

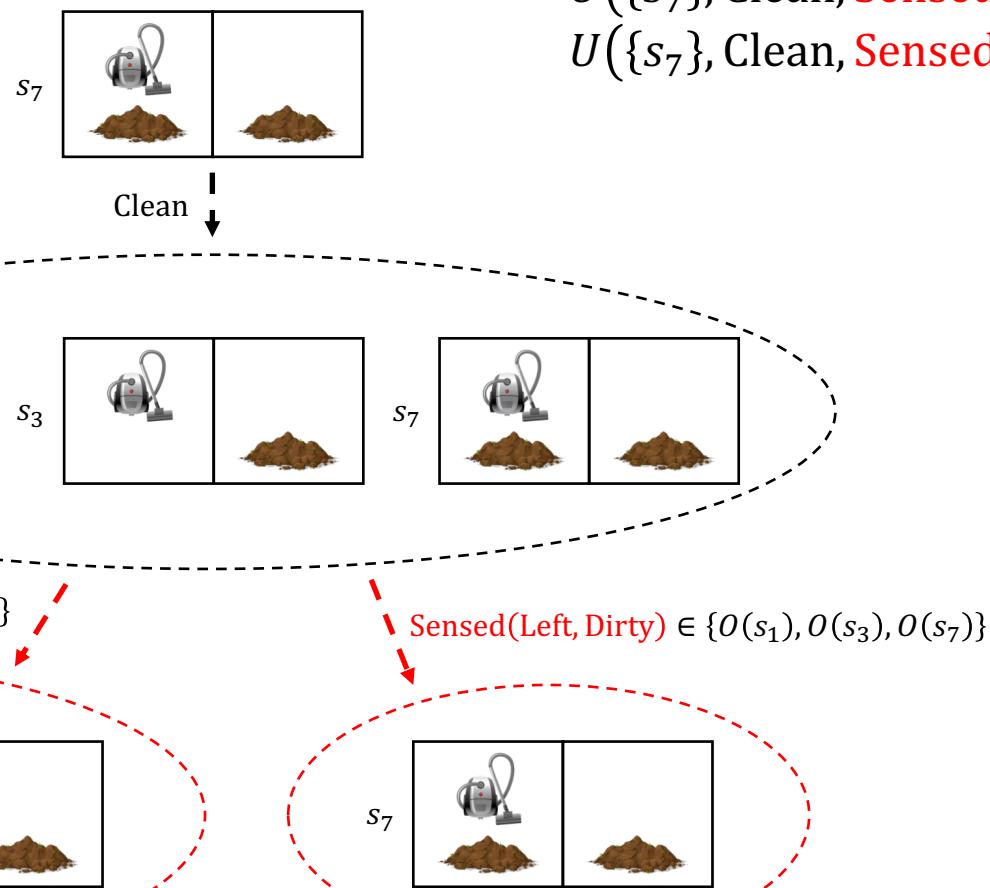
- An update function $U : B \times \mathcal{A} \times \mathcal{P} \rightarrow B$
 - If agent executes action $a \in \mathcal{A}$ in belief state $b \in B$ and observes percept $p \in \mathcal{P}$, then it will be in belief state $U(b, a, p) \in B$
- With a logical model of uncertainty, U would be defined as follows:

$$U(b, a, p) = \{s' \in T(s, a) \mid s \in b, O(s') = p\}$$

- Similar definitions are possible for probabilistic uncertainty, and for noisy sensing...

Reactive Agents

With Belief State: Update – Example



$$U(\{s_7\}, \text{Clean}, \text{Sensed(Left, Clean)}) = \{s_1, s_3\}$$

$$U(\{s_7\}, \text{Clean}, \text{Sensed(Left, Dirty)}) = \{s_7\}$$

Reactive Agents

With Belief State – Example

Algorithm: Rule-based agent program with belief state

Input: 1. Initial belief state b_1
2. Condition-action rule base K

```
1  $b \leftarrow b_1$ 
2 while true do
3   for each  $(\phi, a') \in K$  do
4     if  $b \models \phi$  then
5        $a \leftarrow a'$ 
6       EXECUTE( $a$ )
7       break
8    $p \leftarrow \text{SENSE}()$ 
9    $b \leftarrow \text{UPDATE}(b, a, p)$ 
```

Condition-action rule

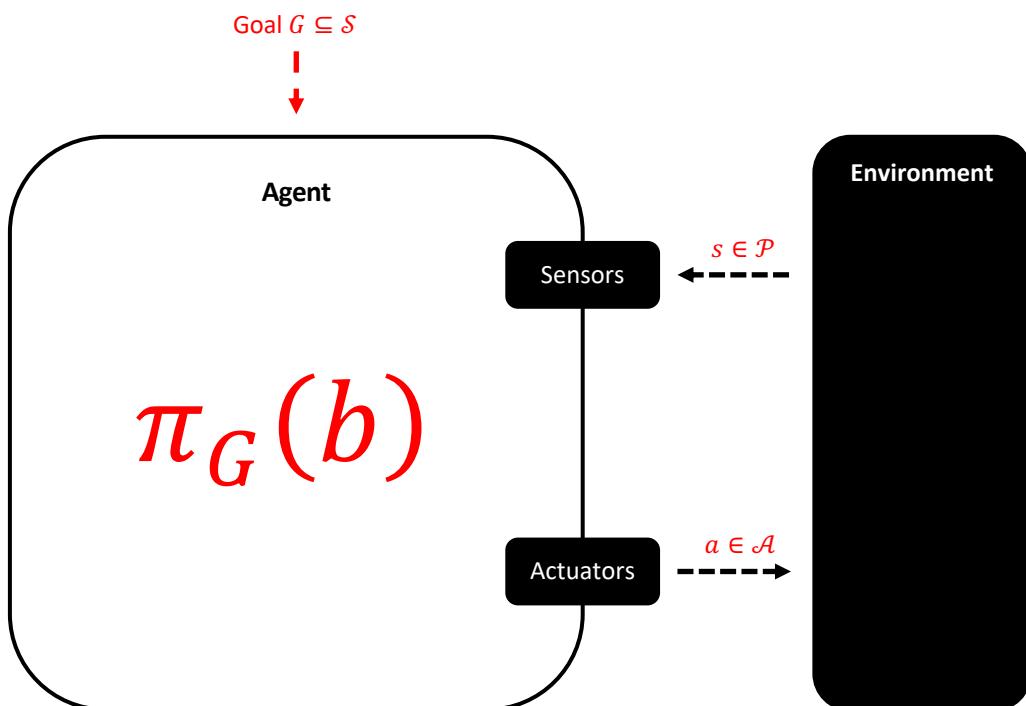
Agent function

#	Condition ϕ	Action $a \in \mathcal{A}$
1	$\text{Agent}(x) \wedge \text{Dirt}(x)$	Clean
2	$\neg \text{Agent}(x) \wedge \text{Dirt}(x)$	Move(x)
3	\top	Null

Belief state $b_i \in B$	Belief-based agent function $\pi_B(b_i)$
$\{s_3, s_4\}$	Move(Right)
$\{s_5, s_6\}$	Move(Left)
$\{s_4, s_5, s_7, s_8\}$	Clean
...	...

Proactive Agents

Goal-Based



- A **goal** is a set of environment states $G \subseteq \mathcal{S}$
 - Goal is **achieved** when the agent arrives in some state $s \in G$
- Agent **terminates** at h if the goal is **believed** to be achieved in belief state associated with h
 - With a **certain** belief state, agent terminates when $b \in G$
 - With a **logically** uncertain belief state, agent terminates when $b \subseteq G$
 - With a **probabilistic** uncertain belief state, agent terminates when e.g. $\sum_{s \in G} b(s) = 1$ or $\sum_{s \in G} b(s) \geq \theta$ with $0 < \theta \leq 1$
- Agent function π_G assigns actions with the objective of **achieving** the goal at some point in the future
 - Agent **continues** to execute actions **until** termination

Proactive Agents

Goal-Based – Example (1)

Algorithm: Goal-based agent program

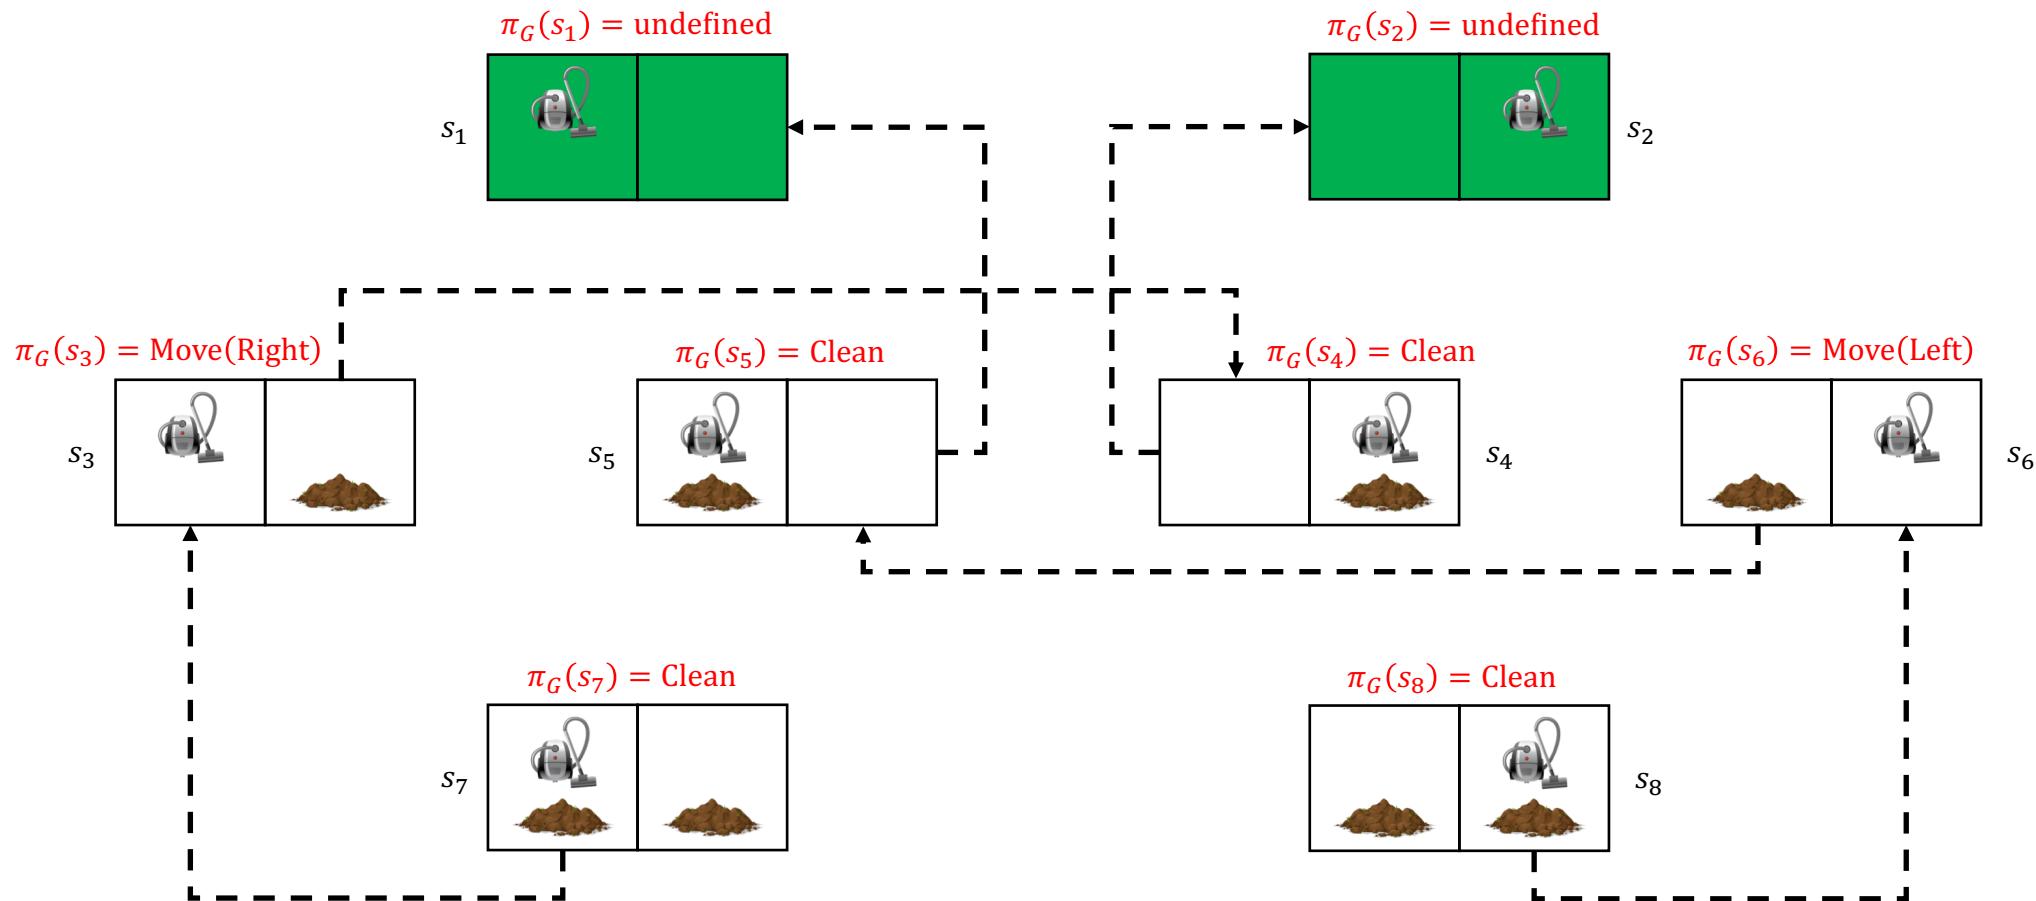
Input: Initial belief state $b_1 \in B$

- 1 $b \leftarrow b_1$
 - 2 $\pi_B \leftarrow \text{PLAN}(b, T, O, G)$ Generate plan for G (subject of next lecture)
 - 3 **while** $b \neq G$ **do**
 - 4 4 $a \leftarrow \pi_B(b)$
 - 5 5 $\text{EXECUTE}(a)$
 - 6 6 $p \leftarrow \text{SENSE}()$
 - 7 7 $b \leftarrow \text{UPDATE}(b, a, p)$
-

Proactive Agents

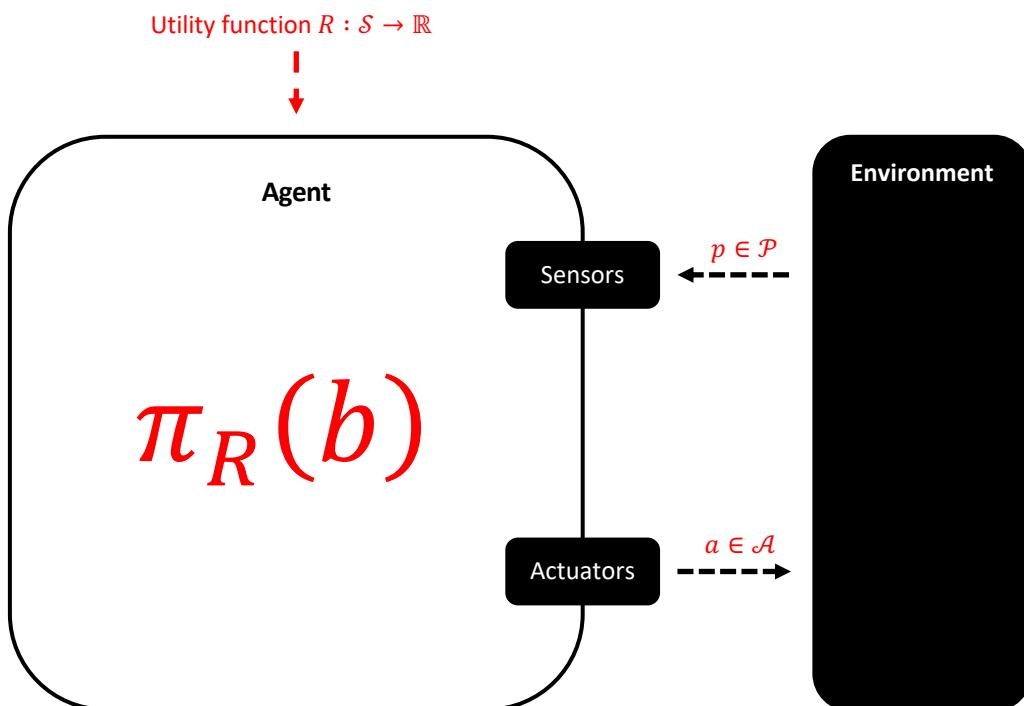
Goal-Based – Example (2)

Goal $G = \{s_1, s_2\}$



Proactive Agents

Utility-Based



- A **utility (or reward) function** is a function $R : \mathcal{S} \rightarrow \mathbb{R}$
 - An environment state $s \in \mathcal{S}$ **preferred** to state $s' \in \mathcal{S}$ if $R(s) > R(s')$
- **Cumulative utility** is utility that accumulates over the agent's life
 - For example, **sum** of utility obtained at all preceding decision-steps
- Agent function π_R assigns actions with the objective of **optimising** cumulative utility over the agent's life
 - Agent **continues** to execute π_U **forever**

Proactive Agents

Utility-Based – Example (1)

Algorithm: Utility-based agent program

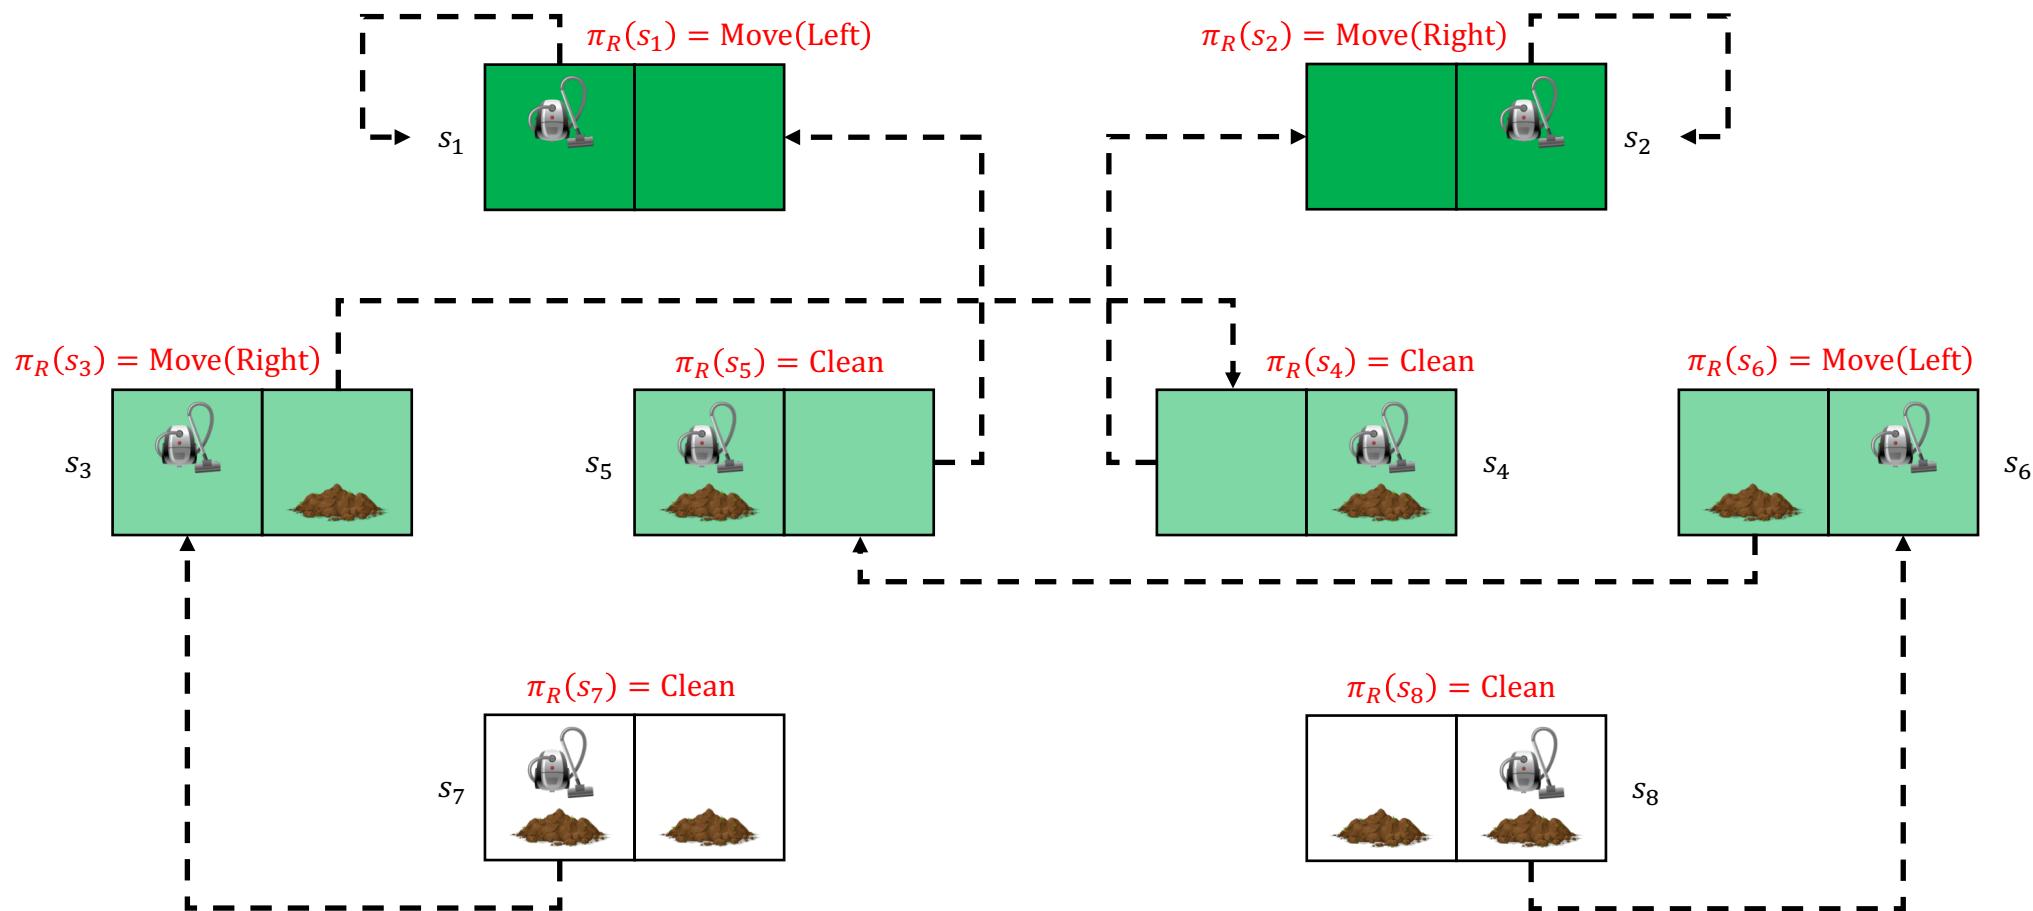
Input: Initial belief state $b_1 \in B$

- 1 $b \leftarrow b_1$
 - 2 $\pi_B \leftarrow \text{PLAN}(b, T, O, R)$ Generate plan for R (subject of next lecture)
 - 3 **while** $true$ **do** Execute plan forever
 - 4 $a \leftarrow \pi_B(b)$
 - 5 EXECUTE(a)
 - 6 $p \leftarrow \text{SENSE}()$
 - 7 $b \leftarrow \text{UPDATE}(b, a, p)$
-

Proactive Agents

Utility-Based – Example (2)

$$\begin{aligned} R(s_1) &= R(s_2) = 2 \\ R(s_3) &= R(s_4) = R(s_5) = R(s_6) = 1 \\ R(s_7) &= R(s_8) = 0 \end{aligned}$$



Cognitive Agents

Motivation

- Need to **balance** reactive and proactive behavior
 - Pure reactive behavior is very **limited** with regards to intelligence and autonomy
 - Pure proactive behavior can be **unnecessary** and/or **unrealistic** (e.g. it is **expensive** to plan)
- **Multi-tasking**
 - Like humans, agents need to pursue **multiple commitments** simultaneously
- **Goal formation**
 - For true autonomy and intelligence, agents need to **form their own goals** through complex reasoning
- **Behavior reconsideration**
 - In **dynamic** and **resource-bounded** environments, agents need to revise their behavior in light of new information
 - For example, agents may need to **drop goals** that are no longer achievable, or may need to **replan**
- **Human-agent interaction**
 - Agents need to be able to **explain** their behavior in a way that is **understandable** to humans
- Many others...

Cognitive Agents

Intentional Systems

- Cognitive agents mirror theories of **human practical reasoning** (especially as proposed by Bratman)
 1. **Deliberation**: deciding **what** states of the world we want to achieve
 2. **Means-end reasoning**: deciding **how** to achieve those states of the world
- Practical reasoning is about figuring out what to do (i.e. reasoning directed **towards action**)
 - Different from **theoretical reasoning** (i.e. reasoning directed **towards beliefs**)
 - *"Practical reasoning is a matter of weighing conflicting considerations for and against competing **options**, where the relevant considerations are provided by what the agent **desires/values/cares about** and what the agent **believes**."* [Bratman]
- Output at each step is a set of **intentions**
 1. Deliberation: **declarative** intentions
 2. Means-end reasoning: **procedural** intentions
- **Planning** (from first-principles) is a form of means-end reasoning
 - So is **hierarchical planning** and **plan selection**

Cognitive Agents

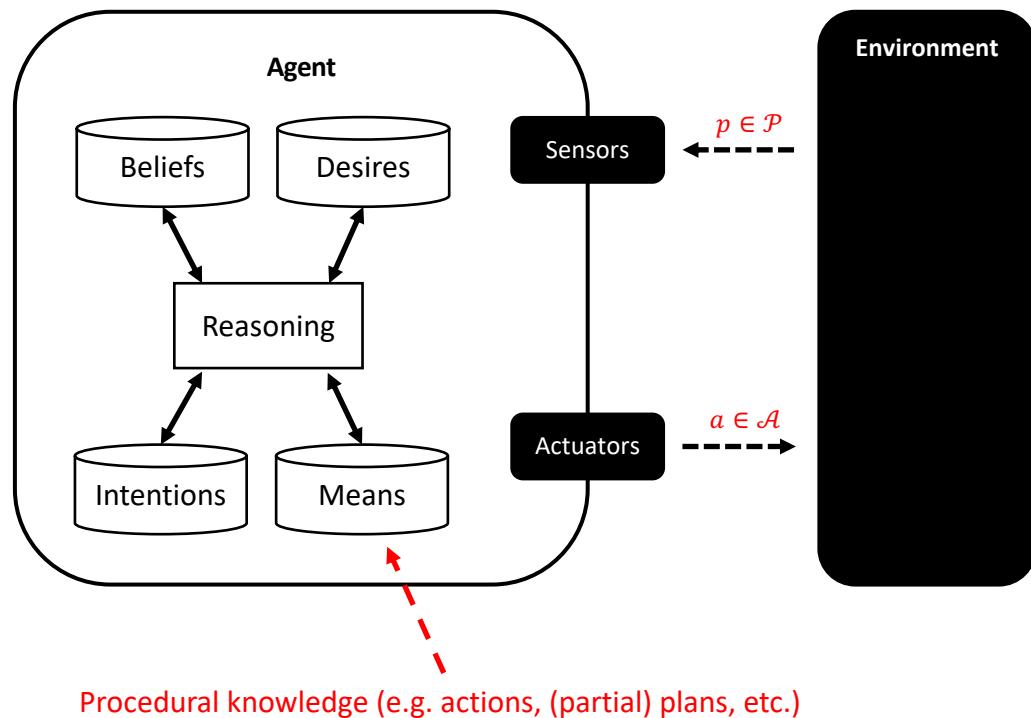
Intentions

Principle	Explanation
Intentions pose problems for agents, who need to determine ways of achieving them	If an agent has an intention φ , you would expect the agent to devote resources to deciding how to bring about φ
Intentions provide a filter for adopting other intentions, which must not conflict	If an agent has an intention φ , you would not expect the agent to adopt an intention ψ that was incompatible with φ
Agents track the success of their intentions, and are inclined to try again if their attempts fail	If an agent's first attempt to achieve φ fails, then all other things being equal, the agent will try an alternative means of achieving φ
Agents believe their intentions are possible	An agent should believe there is at least some way that their intentions could be achieved
Agents do not believe they will not bring about their intentions	It would not be rational for an agent to adopt an intention φ if the agent believed it would fail with φ
Under certain circumstances, agents believe they will bring about their intentions	If the agent intends φ , then the agent believes that under "normal circumstances" it will succeed with φ
Agents need not intend all the expected side effects of their intentions	If the agent believes $\varphi \Rightarrow \psi$ and intends φ , the agent does not necessarily intend ψ also

"My desire to play basketball this afternoon is merely a potential influencer of my conduct this afternoon. It must vie with my other relevant desires [...] before it is settled what I will do. In contrast, once I intend to play basketball this afternoon, the matter is settled: I normally need not continue to weigh the pros and cons. When the afternoon arrives, I will normally just proceed to execute my intentions." [Bratman]

Cognitive Agents

Belief-Desire-Intention (BDI)



- **Beliefs** are the same as **belief states**
 - Sometimes we use the term **belief base** instead (but a belief base is usually represented in some **logical language**)
- **Desires** are states of the world that the agent **wants** to bring about
 - Desires are **like goals** but **do not imply commitment**
 - Desires may be **mutually inconsistent** (i.e. it need not be the case that all desires can be achieved simultaneously)
- **Intentions** are desires that the agent has **committed** to achieving (and possibly the means to do so)
 - See **declarative** vs. **procedural** intentions
 - Declarative intentions are **like goals** in that they **imply commitment**
 - Intentions must be **mutually consistent** and must be **consistent with beliefs**

Cognitive Agents

Belief-Desire-Intention (BDI) – Example

Algorithm: BDI agent program [Wooldridge, 2009]

Input: 1. Initial belief state $b_1 \in B$
2. Initial intention set I_1

```
1  $b \leftarrow b_1$ 
2  $I \leftarrow I_1$ 
3 while true do
4    $D \leftarrow \text{OPTIONS}(b, I)$ 
5    $I \leftarrow \text{FILTER}(b, D, I)$ 
6    $\pi_B \leftarrow \text{PLAN}(b, I)$ 
7   while  $\neg \text{SUCCEEDED}(b, I) \wedge \neg \text{IMPOSSIBLE}(b, I)$  do
8      $a \leftarrow \pi_B(b)$ 
9     EXECUTE( $a$ )
10     $p \leftarrow \text{SENSE}()$ 
11     $b \leftarrow \text{UPDATE}(b, a, p)$ 
12    if RECONSIDER( $b, I$ ) then
13       $D \leftarrow \text{OPTIONS}(b, I)$ 
14       $I \leftarrow \text{FILTER}(b, D, I)$ 
15    if  $\neg \text{SOUND}(\pi_B, b, I)$  then
16       $\pi_B \leftarrow \text{PLAN}(b, I)$ 
```

Deliberation

Means-end reasoning

Generate plan for I

While intentions are believed to be achievable

While intentions are not believed to have been achieved

Decide whether to revise intentions

If (possibly revised) intentions cannot be achieved with current plan

Replan

Intelligent Agents

Some Reflections

- Agent functions are **mathematically identical to plans** (or policies) in automated planning
 - But an agent function is **much more than a plan...**
- Properties of the task environment have **implications** for agent design
 - These properties typically refer to **designer-imposed restrictions** on the model (i.e. **approximations**) rather than inherent properties
- Intelligent agents should **balance** reactive and proactive behavior
 - Proactive behavior can be **slow/expensive** and is **not always necessary**
 - **Humans** behave both proactively and reactively
- Intelligent agents should go **beyond** simple goal- or utility-based agents
 - Agents need to **formulate goals, revise commitments, etc.**
 - BDI is just one approach, but an **influential** one (e.g. see BDI **theory, BDI logics, BDI programming languages**)

This Lecture...

1. Abstract agent representations

- Agent function
- Agent program

2. Task environments

- Problem Analysis
- Properties

3. Types of agents

- Reactive agents
- **Proactive agents**
- Cognitive agents

Next Lecture...

1. Classical planning

- Planning languages

2. Planning under uncertainty

- Conformant planning
- Contingent planning
- Markov decision processes

3. Online planning

- Classical replanning

Bibliography

- Michael Wooldridge. *An Introduction to Multiagent Systems*. John Wiley & Sons, 2nd edition, 2009.

Suggested reading (Chapter 2)
- Stuart J. Russell & Peter Norvig. [*Artificial Intelligence: A Modern Approach*](#). Prentice Hall, 3rd edition, 2009.
- Hector Geffner & Blai Bonet. *A Concise Introduction to Models and Methods for Automated Planning*. Morgan & Claypool, 2013.
- Anand S. Rao. [*AgentSpeak\(L\): BDI agents speak out in a logical computable language*](#). In *Proceedings of the 7th European Workshop on Modelling Autonomous Agents in a Multi-Agent World (MAAMAW'96)*, pages 42-55, 1996.
- Rafael H. Bordini, Jomi Fred Hübner, & Michael Wooldridge. [*Programming Multi-Agent Systems in AgentSpeak Using Jason*](#). John Wiley & Sons, 2007.